

# Quiz 2B

## CS 119: Big Data

Submitted by: Hezekiah Branch

```
In [1]: import pyspark
import nltk
import pandas as pd
```

### Question 1: TF-IDF analysis of text

Read in dataset of last ten president speeches from public Google Cloud bucket

```
In [2]: path = "https://storage.googleapis.com/jsingh-bigdata-public/ten_us_presidents_transcripts.csv"
pres_speeches = pd.read_csv(path)
pres_speeches.drop('Unnamed: 0', axis=1, inplace=True)
pres_speeches.head()
```

```
Out[2]:
```

	content	president	year
0	\nSenator Hatfield, Mr. Chief Justice, Mr. Pre...	Ronald Reagan	1981
1	\nSenator Mathias, Chief Justice Burger, Vice ...	Ronald Reagan	1985
2	\nMr. Chief Justice, Mr. President, Vice Presi...	George Bush	1989
3	\nMy fellow citizens, today we celebrate the m...	William J. Clinton	1993
4	\nMy fellow citizens, at this last Presidentia...	William J. Clinton	1997

```
In [3]: speeches = pres_speeches.copy()
```

```
In [4]: separate_docs = lambda x: x.split('\n')[1:]
speeches['content'] = speeches['content'].apply(separate_docs)
speeches
```

```
Out [4]:
```

	content	president	year
0	[Senator Hatfield, Mr. Chief Justice, Mr. Pres...	Ronald Reagan	1981
1	[Senator Mathias, Chief Justice Burger, Vice P...	Ronald Reagan	1985
2	[Mr. Chief Justice, Mr. President, Vice Presid...	George Bush	1989
3	[My fellow citizens, today we celebrate the my...	William J. Clinton	1993
4	[My fellow citizens, at this last Presidential...	William J. Clinton	1997
5	[Thank you, all. Chief Justice Rehnquist, Pres...	George W. Bush	2001
6	[Vice President Cheney, Mr. Chief Justice, Pre...	George W. Bush	2005
7	[My fellow citizens, I stand here today humble...	Barack Obama	2009
8	[Thank you. Thank you so much., Vice President...	Barack Obama	2013
9	[Chief Justice Roberts, President Carter, Pres...	Donald J. Trump	2017

Add required preprocessing with regex from assignment prompt

```
In [5]: import re
import string

def clean_text(text):
    text = text.lower()
    text = re.sub('[.*?\\]', '', text)
    text = re.sub('[%s]' % re.escape(string.punctuation), '', text)
    text = re.sub('[\d\\n]', '', text)
    return text

round1 = lambda x: clean_text(x).strip()
```

```
In [6]: for row in range(len(speeches.content)):
        cleaned = []
        for item in speeches.content.iloc[row]:
            cleaned.append(clean_text(item).strip())
        speeches.at[row, 'content'] = cleaned
```

```
In [7]: speeches
```

```
Out[7]:
```

	content	president	year
0	[senator hatfield mr chief justice mr pres...	Ronald Reagan	1981
1	[senator mathias chief justice burger vice p...	Ronald Reagan	1985
2	[mr chief justice mr president vice presid...	George Bush	1989
3	[my fellow citizens today we celebrate the my...	William J. Clinton	1993
4	[my fellow citizens at this last presidential...	William J. Clinton	1997
5	[thank you all chief justice rehnquist pres...	George W. Bush	2001
6	[vice president cheney mr chief justice pre...	George W. Bush	2005
7	[my fellow citizens i stand here today humble...	Barack Obama	2009
8	[thank you thank you so much, vice president ...	Barack Obama	2013
9	[chief justice roberts president carter pres...	Donald J. Trump	2017

Build TF-IDF Matrix using scikit-learn TfidfVectorizer

```
In [8]: from sklearn.feature_extraction.text import TfidfVectorizer

corpus = speeches.content.iloc[0]
vectorizer = TfidfVectorizer(analyzer='word', use_idf=True)
tfidf_matrix = vectorizer.fit_transform(corpus)
tfidf_matrix
```

```
Out[8]: <38x835 sparse matrix of type '<class 'numpy.float64'>'
        with 1728 stored elements in Compressed Sparse Row format>
```

```
In [9]: tfidf_matrix[0,:].nonzero()[1]
```

```
Out[9]: array([111, 245, 474,  27, 465, 609, 484, 672,  47, 460,  85, 778, 551,
        388, 105, 470, 313, 635], dtype=int32)
```

```
In [10]: mapping = pd.DataFrame.from_dict(dict(zip(vectorizer.get_feature_names(), vectorizer.idf_)),
        orient='index',
        columns=['tf_idf'])

mapping
```

```
Out[10]:
```

	tf_idf
<b>above</b>	3.970414
<b>abraham</b>	3.970414
<b>accept</b>	3.970414
<b>achieved</b>	3.970414
<b>achievement</b>	3.970414
...	...
<b>yet</b>	3.564949
<b>you</b>	2.584120
<b>young</b>	3.564949
<b>your</b>	3.277267
<b>yourselves</b>	3.970414

835 rows × 1 columns

```
In [11]: top_20 = mapping.tf_idf.sort_values(ascending=False)[0:20]
cell_top_20 = []
for name in top_20.index:
    cell_top_20.append((name, top_20[name]))
cell_top_20
```

```
Out[11]: [('above', 3.970414465569701),
          ('paraphrase', 3.970414465569701),
          ('opening', 3.970414465569701),
          ('opportunities', 3.970414465569701),
          ('order', 3.970414465569701),
          ('orderly', 3.970414465569701),
          ('pace', 3.970414465569701),
          ('paddies', 3.970414465569701),
          ('paid', 3.970414465569701),
          ('parallel', 3.970414465569701),
          ('part', 3.970414465569701),
          ('omaha', 3.970414465569701),
          ('party', 3.970414465569701),
          ('past', 3.970414465569701),
          ('patriotism', 3.970414465569701),
          ('patrol', 3.970414465569701),
          ('perform', 3.970414465569701),
          ('period', 3.970414465569701),
          ('personal', 3.970414465569701),
          ('piled', 3.970414465569701)]
```

Run for entire corpus of presidential speech documents

```
In [12]: tf_idf_matrix = []

for doc in speeches.content:
    corpus = doc
    vectorizer = TfidfVectorizer(analyzer='word', use_idf=True)
    tfidf_matrix = vectorizer.fit_transform(corpus)
    tfidf_matrix
    mapping = pd.DataFrame.from_dict(dict(zip(vectorizer.get_feature_names(), vectorizer.idf_)),
                                     orient='index',
                                     columns=['tf_idf']))

    top_20 = mapping.tf_idf.sort_values(ascending=False)[0:20]
    cell_top_20 = []
    for name in top_20.index:
        cell_top_20.append((name, top_20[name]))
    tf_idf_matrix.append(cell_top_20)
```

In [13]:

```
cols = [str(speeches.iloc[row].president.split()[-1]) + '_' + str(speeches.iloc[row].year) for row in range(10)]
cols = dict(zip(range(10), cols))
tf_idf_matrix = pd.DataFrame(tf_idf_matrix).transpose()
tf_idf_matrix = tf_idf_matrix.rename(columns=cols)
tf_idf_matrix
```

Out[13]:

	Reagan_1981	Reagan_1985	Bush_1989	Clinton_1993	Clinton_1997	Bush_2001
0	(above, 3.970414465569701)	(absent, 4.091042453358316)	(important, 3.70805020110221)	(abiding, 3.2512917986064953)	(abolished, 3.5649493574615367)	(abandonment, 3.772588722239781)
1	(paraphrase, 3.970414465569701)	(out, 4.091042453358316)	(money, 3.70805020110221)	(ocean, 3.2512917986064953)	(opened, 3.5649493574615367)	(others, 3.772588722239781)
2	(opening, 3.970414465569701)	(once, 4.091042453358316)	(merely, 3.70805020110221)	(part, 3.2512917986064953)	(obligations, 3.5649493574615367)	(persistent, 3.772588722239781)
3	(opportunities, 3.970414465569701)	(open, 4.091042453358316)	(michel, 3.70805020110221)	(pain, 3.2512917986064953)	(obsessions, 3.5649493574615367)	(permit, 3.772588722239781)
4	(order, 3.970414465569701)	(opportunities, 4.091042453358316)	(mistrust, 3.70805020110221)	(over, 3.2512917986064953)	(off, 3.5649493574615367)	(people, 3.772588722239781)
5	(orderly, 3.970414465569701)	(oppression, 4.091042453358316)	(mists, 3.70805020110221)	(out, 3.2512917986064953)	(office, 3.5649493574615367)	(peaceful, 3.772588722239781)
6	(pace, 3.970414465569701)	(orderly, 4.091042453358316)	(mitchell, 3.70805020110221)	(order, 3.2512917986064953)	(onto, 3.5649493574615367)	(peace, 3.772588722239781)
7	(paddies, 3.970414465569701)	(origin, 4.091042453358316)	(trumpets, 3.70805020110221)	(opportunities, 3.2512917986064953)	(open, 3.5649493574615367)	(pastor, 3.772588722239781)
8	(paid, 3.970414465569701)	(others, 4.091042453358316)	(truly, 3.70805020110221)	(one, 3.2512917986064953)	(opportunities, 3.5649493574615367)	(passing, 3.772588722239781)
9	(parallel, 3.970414465569701)	(ours, 4.091042453358316)	(member, 3.70805020110221)	(oldest, 3.2512917986064953)	(nuclear, 3.5649493574615367)	(passed, 3.772588722239781)
10	(part, 3.970414465569701)	(overwhelm, 4.091042453358316)	(most, 3.70805020110221)	(older, 3.2512917986064953)	(other, 3.5649493574615367)	(pain, 3.772588722239781)
11	(omaha, 3.970414465569701)	(planter, 4.091042453358316)	(motives, 3.70805020110221)	(old, 3.2512917986064953)	(others, 3.5649493574615367)	(page, 3.772588722239781)
12	(party, 3.970414465569701)	(own, 4.091042453358316)	(move, 3.70805020110221)	(often, 3.2512917986064953)	(outlines, 3.5649493574615367)	(over, 3.772588722239781)

13	(past, 3.970414465569701)	(paces, 4.091042453358316)	(moves, 3.70805020110221)	(office, 3.2512917986064953)	(overcome, 3.5649493574615367)	(out, 3.772588722239781)
14	(patriotism, 3.970414465569701)	(part, 4.091042453358316)	(moving, 3.70805020110221)	(offer, 3.2512917986064953)	(parents, 3.5649493574615367)	(order, 3.772588722239781)
15	(patrol, 3.970414465569701)	(parties, 4.091042453358316)	(triumph, 3.70805020110221)	(oath, 3.2512917986064953)	(part, 3.5649493574615367)	(remains, 3.772588722239781)
16	(perform, 3.970414465569701)	(party, 4.091042453358316)	(try, 3.70805020110221)	(pays, 3.2512917986064953)	(numbers, 3.5649493574615367)	(options, 3.772588722239781)
17	(period, 3.970414465569701)	(pass, 4.091042453358316)	(meet, 3.70805020110221)	(numbers, 3.2512917986064953)	(nothing, 3.5649493574615367)	(opportunity, 3.772588722239781)
18	(personal, 3.970414465569701)	(passing, 4.091042453358316)	(voices, 3.70805020110221)	(nor, 3.2512917986064953)	(partisanship, 3.5649493574615367)	(onward, 3.772588722239781)
19	(piled, 3.970414465569701)	(paying, 4.091042453358316)	(turning, 3.70805020110221)	(news, 3.2512917986064953)	(mystical, 3.5649493574615367)	(one, 3.772588722239781)

## Question 1: Retail Data Analysis

```
In [14]: online_retail = pd.read_csv("online-retail-online_retail_II.csv")
online_retail
```

Out[14]:

	Invoice	StockCode	Description	Quantity	InvoiceDate	Price	Customer ID	Country
0	489434	85048	15CM CHRISTMAS GLASS BALL 20 LIGHTS	12	12/1/2009 7:45	6.95	13085.0	United Kingdom
1	489434	79323P	PINK CHERRY LIGHTS	12	12/1/2009 7:45	6.75	13085.0	United Kingdom
2	489434	79323W	WHITE CHERRY LIGHTS	12	12/1/2009 7:45	6.75	13085.0	United Kingdom
3	489434	22041	RECORD FRAME 7" SINGLE SIZE	48	12/1/2009 7:45	2.10	13085.0	United Kingdom
4	489434	21232	STRAWBERRY CERAMIC TRINKET BOX	24	12/1/2009 7:45	1.25	13085.0	United Kingdom
...	...	...	...	...	...	...	...	...
1067366	581587	22899	CHILDREN'S APRON DOLLY GIRL	6	12/9/2011 12:50	2.10	12680.0	France
1067367	581587	23254	CHILDRENS CUTLERY DOLLY GIRL	4	12/9/2011 12:50	4.15	12680.0	France
1067368	581587	23255	CHILDRENS CUTLERY CIRCUS PARADE	4	12/9/2011 12:50	4.15	12680.0	France
1067369	581587	22138	BAKING SET 9 PIECE RETROSPOT	3	12/9/2011 12:50	4.95	12680.0	France
1067370	581587	POST	POSTAGE	1	12/9/2011 12:50	18.00	12680.0	France

1067371 rows x 8 columns

In [15]:

```
from pyspark.sql import SparkSession

spark = SparkSession \
    .builder \
    .appName("Online Retail") \
    .getOrCreate()
```



```
21/12/29 23:03:34 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using b
uiltin-java classes where applicable
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
```

In [16]:

```
path = "online-retail-online_retail_II.csv"

df = spark.read.option("header", True).csv(path, inferSchema=True)
df.show()
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+
|Invoice|StockCode|Description|Quantity|InvoiceDate|Price|Customer ID|Country|
+-----+-----+-----+-----+-----+-----+-----+-----+
| 489434| 85048|15CM CHRISTMAS GL...| 12|12/1/2009 7:45| 6.95| 13085|United Kingdom|
| 489434| 79323P| PINK CHERRY LIGHTS| 12|12/1/2009 7:45| 6.75| 13085|United Kingdom|
| 489434| 79323W| WHITE CHERRY LIGHTS| 12|12/1/2009 7:45| 6.75| 13085|United Kingdom|
| 489434| 22041|"RECORD FRAME 7"...| 48|12/1/2009 7:45| 2.1| 13085|United Kingdom|
| 489434| 21232|STRAWBERRY CERAMI...| 24|12/1/2009 7:45| 1.25| 13085|United Kingdom|
| 489434| 22064|PINK DOUGHNUT TRI...| 24|12/1/2009 7:45| 1.65| 13085|United Kingdom|
| 489434| 21871| SAVE THE PLANET MUG| 24|12/1/2009 7:45| 1.25| 13085|United Kingdom|
| 489434| 21523|FANCY FONT HOME S...| 10|12/1/2009 7:45| 5.95| 13085|United Kingdom|
| 489435| 22350| CAT BOWL| 12|12/1/2009 7:46| 2.55| 13085|United Kingdom|
| 489435| 22349|DOG BOWL , CHASIN...| 12|12/1/2009 7:46| 3.75| 13085|United Kingdom|
| 489435| 22195|HEART MEASURING S...| 24|12/1/2009 7:46| 1.65| 13085|United Kingdom|
| 489435| 22353|LUNCHBOX WITH CUT...| 12|12/1/2009 7:46| 2.55| 13085|United Kingdom|
| 489436| 48173C|DOOR MAT BLACK FL...| 10|12/1/2009 9:06| 5.95| 13078|United Kingdom|
| 489436| 21755|LOVE BUILDING BLO...| 18|12/1/2009 9:06| 5.45| 13078|United Kingdom|
| 489436| 21754|HOME BUILDING BLO...| 3|12/1/2009 9:06| 5.95| 13078|United Kingdom|
| 489436| 84879|ASSORTED COLOUR B...| 16|12/1/2009 9:06| 1.69| 13078|United Kingdom|
| 489436| 22119| PEACE WOODEN BLO...| 3|12/1/2009 9:06| 6.95| 13078|United Kingdom|
| 489436| 22142|CHRISTMAS CRAFT W...| 12|12/1/2009 9:06| 1.45| 13078|United Kingdom|
| 489436| 22296|HEART IVORY TRELL...| 12|12/1/2009 9:06| 1.65| 13078|United Kingdom|
| 489436| 22295|HEART FILIGREE DO...| 12|12/1/2009 9:06| 1.65| 13078|United Kingdom|
+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 20 rows
```

In [17]:

```
df.dtypes
```

```
Out[17]: [('Invoice', 'string'),
          ('StockCode', 'string'),
          ('Description', 'string'),
          ('Quantity', 'int'),
          ('InvoiceDate', 'string'),
          ('Price', 'double'),
          ('Customer ID', 'int'),
          ('Country', 'string')]
```

Question: Loading the data into a Dataframe and removing junk records. How many records were removed by doing so?

Response: **When removing junk records, we have found 269486 records to remove.**

```
In [18]: df_clean = df.dropna()
df_clean = df_clean.dropDuplicates()
num_empty_rows = df.count() - df_clean.count()
num_empty_rows
```

```
Out[18]: 269486
```

Dealing with calculation of monetary value

```
In [19]: df_clean.show()
```

[Stage 8:=====>

(1 + 3) / 4]

Invoice	StockCode	Description	Quantity	InvoiceDate	Price	Customer ID	Country
489514	21791	VINTAGE HEADS AND...	72	12/1/2009 11:21	1.06	15311	United Kingdom
489520	72739B	WHITE CHOCOLATE S...	12	12/1/2009 11:41	1.25	14911	EIRE
489522	22315	200 RED + WHITE B...	1	12/1/2009 11:45	1.25	15998	United Kingdom
489536	21611	SET OF 12 LILY BO...	2	12/1/2009 12:13	2.95	16393	United Kingdom
489561	21816	CHRISTMAS TREE T-...	6	12/1/2009 12:57	1.45	14654	United Kingdom
489562	35071	ASSORTED SANTA CH...	1	12/1/2009 13:07	0.85	17998	United Kingdom
489576	22152	PLACE SETTING WHI...	24	12/1/2009 13:38	0.42	15984	United Kingdom
489599	21239	PINK SPOTTY CUP	16	12/1/2009 14:40	0.85	12758	Portugal
489658	79323LP	LIGHT PINK CHERRY...	6	12/1/2009 17:31	6.75	15485	United Kingdom
489679	22086	PAPER CHAIN KIT 5...	6	12/2/2009 10:00	2.95	16163	United Kingdom
489681	85226C	BLUE PULL BACK RA...	11	12/2/2009 10:02	0.55	17998	United Kingdom
489723	85231E	STRAWBERRY SCENTE...	36	12/2/2009 10:58	0.85	14299	United Kingdom
489765	22315	200 RED + WHITE B...	1	12/2/2009 11:33	1.25	15353	United Kingdom
489766	21975	PACK OF 60 DINOSA...	1	12/2/2009 11:34	0.55	17818	United Kingdom
489791	84029G	KNITTED UNION FLA...	1	12/2/2009 12:06	3.75	15542	United Kingdom
489797	72807A	SET/3 ROSE CANDLE...	5	12/2/2009 12:19	4.25	15581	United Kingdom
489814	21973	SET OF 36 MUSHROO...	2	12/2/2009 13:06	1.45	14669	United Kingdom
489827	22335	HEART DECORATION ...	24	12/2/2009 13:51	0.65	17412	United Kingdom
489866	21877	HOME SWEET HOME MUG	6	12/2/2009 15:04	1.25	16200	United Kingdom
489878	20966	SANDWICH BATH SPONGE	1	12/2/2009 15:51	1.25	15989	United Kingdom

only showing top 20 rows

Change the Price and InvoiceDate columns to type Double and timestamp, respectively.

The cell ran but I cleared the output due to a large number of observations

In [20]:

```
from pyspark.sql.types import DoubleType, TimestampType, StringType, DateType

df_clean = df_clean.withColumn("Price", df_clean["Price"].cast(DoubleType()))
```

```
In [22]: from pyspark.sql.functions import to_timestamp

spark.sql("set spark.sql.legacy.timeParserPolicy=LEGACY")

test = df_clean.select('Customer ID', to_timestamp(df_clean.InvoiceDate, 'MM/d/yyyy HH:mm').alias('InvoiceDate'))
```

```
In [23]: test.schema['InvoiceDate']
```

```
Out[23]: StructField(InvoiceDate,TimestampType,true)
```

```
In [25]: test = test.orderBy('InvoiceDate', ascending=False)
```

```
In [26]: test.show()
```

[Stage 11:=====>

(154 + 4) / 200]

Customer ID	InvoiceDate
12680	2011-12-09 12:50:00
12680	2011-12-09 12:50:00
12680	2011-12-09 12:50:00
12680	2011-12-09 12:50:00
12680	2011-12-09 12:50:00
12680	2011-12-09 12:50:00
12680	2011-12-09 12:50:00
12680	2011-12-09 12:50:00
12680	2011-12-09 12:50:00
12680	2011-12-09 12:50:00
12680	2011-12-09 12:50:00
12680	2011-12-09 12:50:00
12680	2011-12-09 12:50:00
12680	2011-12-09 12:50:00
12680	2011-12-09 12:50:00
12680	2011-12-09 12:50:00
12680	2011-12-09 12:50:00
13113	2011-12-09 12:49:00
13113	2011-12-09 12:49:00
13113	2011-12-09 12:49:00
13113	2011-12-09 12:49:00

only showing top 20 rows

In [27]:

```
test.count()
```

Out[27]:

797885

In [28]:

```
df_clean.count()
```

Out[28]:

797885

Order the Spark DataFrame by Price descending to calculate percentiles

```
In [29]: df_clean.orderBy("Price", ascending=False).show(20)
```

```
[Stage 18:=====>                                     (1 + 3) / 4]
```

Invoice	StockCode	Description	Quantity	InvoiceDate	Price	Customer ID	Country
C556445	M	Manual	-1	6/10/2011 15:31	38970.0	15098	United Kingdom
C512770	M	Manual	-1	6/17/2010 16:52	25111.09	17399	United Kingdom
C502262	M	Manual	-1	3/23/2010 15:20	10953.5	12918	United Kingdom
C502264	M	Manual	-1	3/23/2010 15:24	10953.5	12918	United Kingdom
502263	M	Manual	1	3/23/2010 15:22	10953.5	12918	United Kingdom
C522793	M	Manual	-1	9/16/2010 14:53	10468.8	14063	United Kingdom
524159	M	Manual	1	9/27/2010 16:12	10468.8	14063	United Kingdom
C525398	M	Manual	-1	10/5/2010 11:47	10468.8	14063	United Kingdom
496115	M	Manual	1	1/29/2010 11:04	8985.6	17949	United Kingdom
C496116	M	Manual	-1	1/29/2010 11:05	8985.6	17949	United Kingdom
551697	POST	POSTAGE	1	5/3/2011 13:46	8142.75	16029	United Kingdom
C551685	POST	POSTAGE	-1	5/3/2011 12:51	8142.75	16029	United Kingdom
C525470	M	Manual	-1	10/5/2010 15:12	7044.79	15413	United Kingdom
501768	M	Manual	1	3/19/2010 11:45	6958.17	15760	Norway
C501751	M	Manual	-1	3/19/2010 11:30	6958.17	15760	Norway
C501769	M	Manual	-1	3/19/2010 11:49	6958.17	15760	Norway
501766	M	Manual	1	3/19/2010 11:35	6958.17	15760	Norway
C551699	M	Manual	-1	5/3/2011 14:12	6930.0	16029	United Kingdom
C505490	M	Manual	-1	4/22/2010 12:55	5876.34	15849	United Kingdom
C504637	M	Manual	-1	4/15/2010 14:07	5843.7	17017	United Kingdom

only showing top 20 rows

```
In [30]: monetary = df_clean.groupBy('Customer ID').sum('Price')
monetary.show()
```

```
[Stage 21:=====> (195 + 4) / 200]
```

Customer ID	sum(Price)
15727	2446.9100000000008
16503	883.9200000000009
17753	198.98999999999998
15957	356.14999999999999
16386	285.56999999999994
17389	2438.0299999999997
12940	405.89999999999986
16574	155.11999999999995
13623	1051.0200000000002
13832	148.44
16861	131.48000000000002
13285	539.3000000000001
17679	291.41
17420	177.34000000000003
15619	13.25
15790	117.75000000000001
18051	113.35999999999999
16339	89.25000000000001
14570	223.49999999999986
13840	114.98000000000002

only showing top 20 rows

```
In [31]: from pyspark.sql.functions import percentile_approx

quantiles = monetary.select(percentile_approx("sum(Price)", [0.85, 0.70, 0.40], 1000000).alias("Top 15-30-"))

In [32]: quantiles.show(3, False)
```

```
[Stage 25:=====> (193 + 4) / 200]
```

```
+-----+
|Top 15-30-60th %|
+-----+
|[734.13000000000003, 355.20999999999987, 118.50999999999999]|
+-----+
```

Use percentiles to assign values to each of the customer records based on the provided chart

In [33]:

```
from pyspark.sql.functions import when

df_clean = df_clean.withColumn("Monetary", \
    when((df_clean["Price"] >= 734.13000000000006), 1)
    .when((df_clean["Price"] >= 355.21) & (df_clean["Price"] < 734.13000000000006), 2)
    .when((df_clean["Price"] >= 118.51000000000002) & (df_clean["Price"] < 355.21), 3)
    .otherwise(4))

df_clean.show()
```

```
[Stage 27:=====> (1 + 3) / 4]
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+
|Invoice|StockCode|Description|Quantity|InvoiceDate|Price|Customer ID|Country|Monetary|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+
| 489514| 21791|VINTAGE HEADS AND...| 72|12/1/2009 11:21| 1.06| 15311|United Kingdom| 4
| 489520| 72739B|WHITE CHOCOLATE S...| 12|12/1/2009 11:41| 1.25| 14911| EIRE| 4
| 489522| 22315|200 RED + WHITE B...| 1|12/1/2009 11:45| 1.25| 15998|United Kingdom| 4
| 489536| 21611|SET OF 12 LILY BO...| 2|12/1/2009 12:13| 2.95| 16393|United Kingdom| 4
| 489561| 21816|CHRISTMAS TREE T-...| 6|12/1/2009 12:57| 1.45| 14654|United Kingdom| 4
| 489562| 35071|ASSORTED SANTA CH...| 1|12/1/2009 13:07| 0.85| 17998|United Kingdom| 4
| 489576| 22152|PLACE SETTING WHI...| 24|12/1/2009 13:38| 0.42| 15984|United Kingdom| 4
```



489599	21239	PINK SPOTTY CUP	16	12/1/2009 14:40	0.85	12758	Portugal	4
489658	79323LP	LIGHT PINK CHERRY...	6	12/1/2009 17:31	6.75	15485	United Kingdom	4
489679	22086	PAPER CHAIN KIT 5...	6	12/2/2009 10:00	2.95	16163	United Kingdom	4
489681	85226C	BLUE PULL BACK RA...	11	12/2/2009 10:02	0.55	17998	United Kingdom	4
489723	85231E	STRAWBERRY SCENTE...	36	12/2/2009 10:58	0.85	14299	United Kingdom	4
489765	22315	200 RED + WHITE B...	1	12/2/2009 11:33	1.25	15353	United Kingdom	4
489766	21975	PACK OF 60 DINOSA...	1	12/2/2009 11:34	0.55	17818	United Kingdom	4
489791	84029G	KNITTED UNION FLA...	1	12/2/2009 12:06	3.75	15542	United Kingdom	4
489797	72807A	SET/3 ROSE CANDLE...	5	12/2/2009 12:19	4.25	15581	United Kingdom	4
489814	21973	SET OF 36 MUSHROO...	2	12/2/2009 13:06	1.45	14669	United Kingdom	4
489827	22335	HEART DECORATION ...	24	12/2/2009 13:51	0.65	17412	United Kingdom	4
489866	21877	HOME SWEET HOME MUG	6	12/2/2009 15:04	1.25	16200	United Kingdom	4
489878	20966	SANDWICH BATH SPONGE	1	12/2/2009 15:51	1.25	15989	United Kingdom	4

+-----+-----+-----+-----+-----+-----+-----+-----+-----+

+  
only showing top 20 rows



Repeating steps for frequency

```
In [34]: df_clean.dtypes
```

```
Out[34]: [('Invoice', 'string'),
          ('StockCode', 'string'),
          ('Description', 'string'),
          ('Quantity', 'int'),
          ('InvoiceDate', 'string'),
          ('Price', 'double'),
          ('Customer ID', 'int'),
          ('Country', 'string'),
          ('Monetary', 'int')]
```

```
In [35]: df_clean = df_clean.withColumnRenamed("Customer ID", "Customer_ID")
```

```
In [36]: frequency = df_clean.withColumn("Quantity", df_clean["Quantity"].cast(DoubleType()))
```

```
In [37]: from pyspark.sql.functions import countDistinct

freq_view = frequency.groupBy('Customer_ID').agg(countDistinct('Invoice'))
freq_view.show()
```

```
[Stage 31:=====> (192 + 4) / 200]
```

Customer_ID	count(Invoice)
16574	3
15727	15
17389	77
15619	1
15447	6
18051	8
13623	15
12940	4
14450	7
16503	13
15846	1
14832	3
15790	1
13285	6
17753	5
14570	3
13832	3
17679	11
16861	6
15957	3

only showing top 20 rows

In [38]:

```
freq_quantiles = freq_view.select(percentile_approx("count(Invoice)", [0.85, 0.70, 0.40], 1000000).alias("freq_quantiles.show()"))
```

```
[Stage 36:=====> (185 + 5) / 200]
```

Top 15-30-60th %
[13, 7, 3]

In [39]:

```
freq_view = freq_view.withColumn("Frequency", \
    when((freq_view["count(Invoice)"] >= 13), 1)
    .when((freq_view["count(Invoice)"] >= 7) & (freq_view["count(Invoice)"] < 13), 2)
    .when((freq_view["count(Invoice)"] >= 3) & (freq_view["count(Invoice)"] < 7), 3)
    .otherwise(4))

freq_view.show()
```

[Stage 40:=====> (190 + 4) / 200]

Customer_ID	count(Invoice)	Frequency
16574	3	3
15727	15	1
17389	77	1
15619	1	4
15447	6	3
18051	8	2
13623	15	1
12940	4	3
14450	7	2
16503	13	1
15846	1	4
14832	3	3
15790	1	4
13285	6	3
17753	5	3
14570	3	3
13832	3	3
17679	11	2
16861	6	3
15957	3	3

only showing top 20 rows

In [40]:

```
freq_view = freq_view.withColumnRenamed("Customer_ID", "ID")
freq_view = freq_view.drop("count(Invoice)")
```

```
In [41]: freq_df = df_clean.join(freq_view, df_clean["Customer_ID"] == freq_view["ID"], "left")
```

```
In [42]: freq_df = freq_df.drop("ID")
freq_df.show()
```

```
[Stage 46:=====>(199 + 1) / 200]
```

Invoice	StockCode	Description	Quantity	InvoiceDate	Price	Customer_ID	Country	Monetary
Frequency								
513796	85017B	ENVELOPE 50 BLOSS...	12	6/28/2010 15:57	0.85	12799	Japan	
4	4							
513796	22077	6 RIBBONS RUSTIC ...	12	6/28/2010 15:57	1.65	12799	Japan	
4	4							
513796	22509	SEWING BOX RETROS...	1	6/28/2010 15:57	16.95	12799	Japan	
4	4							
513796	85032C	CURIOUS IMAGES GI...	6	6/28/2010 15:57	2.1	12799	Japan	
4	4							
513796	85049G	CHOCOLATE BOX RIB...	12	6/28/2010 15:57	1.25	12799	Japan	
4	4							
513796	85032A	ROMANTIC IMAGES G...	6	6/28/2010 15:57	2.1	12799	Japan	
4	4							
513796	22074	6 RIBBONS SHIMMER...	12	6/28/2010 15:57	1.65	12799	Japan	
4	4							
513796	85032B	BLOSSOM IMAGES GI...	6	6/28/2010 15:57	2.1	12799	Japan	
4	4							
513796	22078	RIBBON REEL LACE ...	10	6/28/2010 15:57	2.1	12799	Japan	
4	4							
513796	21259	VICTORIAN SEWING ...	2	6/28/2010 15:57	5.95	12799	Japan	
4	4							
513796	85017C	ENVELOPE 50 CURIO...	12	6/28/2010 15:57	0.85	12799	Japan	
4	4							
513796	85178	VICTORIAN SEWING KIT	12	6/28/2010 15:57	1.25	12799	Japan	
4	4							
513796	85176	SEWING SUSAN 21 N...	12	6/28/2010 15:57	0.85	12799	Japan	
4	4							
513796	85049D	BRIGHT BLUES RIBB...	12	6/28/2010 15:57	1.25	12799	Japan	
4	4							

513796	22081	RIBBON REEL FLORA...		10	6/28/2010 15:57	1.65	12799	Japan
4	4							
566488	22600	CHRISTMAS RETROSP...		12	9/13/2011 10:16	0.85	12940	United Kingdom
4	3							
571270	22696	WICKER WREATH LARGE		2	10/16/2011 12:09	1.95	12940	United Kingdom
4	3							
571270	21619	4 VANILLA BOTANIC...		3	10/16/2011 12:09	1.25	12940	United Kingdom
4	3							
566488	23174	REGENCY SUGAR BOW...		4	9/13/2011 10:16	4.15	12940	United Kingdom
4	3							
571270	23333	IVORY WICKER HEAR...		4	10/16/2011 12:09	1.25	12940	United Kingdom
4	3							

+-----+-----+-----+-----+-----+-----+-----+-----+-----+  
 -+-----+  
 only showing top 20 rows



Moving on to the last part of segmentation on recency

In [43]:

```

from pyspark.sql.functions import first

dated = test.groupBy('Customer ID').agg(first('InvoiceDate'))

dated.show()

```

```
[Stage 52:=====> (190 + 4) / 200]
```

```
+-----+-----+
|Customer ID| first(InvoiceDate)|
+-----+-----+
|      17389|2011-12-09 09:38:00|
|      15790|2011-11-29 14:53:00|
|      15619|2011-11-29 08:14:00|
|      15727|2011-11-23 12:36:00|
|      13832|2011-11-22 12:31:00|
|      13285|2011-11-16 13:19:00|
|      16386|2011-11-11 12:28:00|
|      13623|2011-11-09 12:00:00|
|      15957|2011-11-08 12:14:00|
|      12940|2011-10-24 14:04:00|
|      17420|2011-10-20 14:52:00|
|      17679|2011-10-18 07:43:00|
|      16861|2011-10-11 09:05:00|
|      16574|2011-09-29 13:39:00|
|      16503|2011-08-25 11:46:00|
|      18024|2011-07-10 12:40:00|
|      14450|2011-06-12 10:46:00|
|      14570|2011-03-04 10:58:00|
|      16339|2011-02-28 13:41:00|
|      15447|2011-01-13 11:26:00|
+-----+-----+
```

only showing top 20 rows

In [44]:

```
recency_view = dated.withColumn("Recency", \
    when((dated["first(InvoiceDate)"] >= "11/15/2011"), 1)
    .when((dated["first(InvoiceDate)"] >= "9/4/2011") & (dated["first(InvoiceDate)"] < "11/14/2011"), 2)
    .when((dated["first(InvoiceDate)"] >= "1/4/2011") & (dated["first(InvoiceDate)"] < "9/4/2011"), 3)
    .otherwise(4))

recency_view.show()
```

```
[Stage 58:=====> (190 + 4) / 200]
```

Customer ID	first(InvoiceDate)	Recency
17389	2011-12-09 09:38:00	4
15790	2011-11-29 14:53:00	4
15619	2011-11-29 08:14:00	4
15727	2011-11-23 12:36:00	4
13832	2011-11-22 12:31:00	4
13285	2011-11-16 13:19:00	4
16386	2011-11-11 12:28:00	4
13623	2011-11-09 12:00:00	4
15957	2011-11-08 12:14:00	4
12940	2011-10-24 14:04:00	4
17420	2011-10-20 14:52:00	4
17679	2011-10-18 07:43:00	4
16861	2011-10-11 09:05:00	4
16574	2011-09-29 13:39:00	4
16503	2011-08-25 11:46:00	4
18024	2011-07-10 12:40:00	4
14450	2011-06-12 10:46:00	4
14570	2011-03-04 10:58:00	4
16339	2011-02-28 13:41:00	4
15447	2011-01-13 11:26:00	4

only showing top 20 rows

In [45]:

```
recency_df = freq_df.join(recency_view, freq_df["Customer_ID"] == recency_view["Customer ID"], "left")
recency_df.show()
```

```
[Stage 68:=====> (195 + 4) / 200]
```

Invoice	StockCode	Description	Quantity	InvoiceDate	Price	Customer_ID	Country	Monetary
513796	85017B	ENVELOPE 50 BLOSS...	12	6/28/2010 15:57	0.85	12799	Japan	
4	4	12799	2010-06-28 15:57:00	4				



513796	22077	6 RIBBONS RUSTIC ...	12	6/28/2010 15:57	1.65	12799	Japan
4	4	12799 2010-06-28 15:57:00	4				
513796	22509	SEWING BOX RETROS...	1	6/28/2010 15:57	16.95	12799	Japan
4	4	12799 2010-06-28 15:57:00	4				
513796	85032C	CURIOUS IMAGES GI...	6	6/28/2010 15:57	2.1	12799	Japan
4	4	12799 2010-06-28 15:57:00	4				
513796	85049G	CHOCOLATE BOX RIB...	12	6/28/2010 15:57	1.25	12799	Japan
4	4	12799 2010-06-28 15:57:00	4				
513796	85032A	ROMANTIC IMAGES G...	6	6/28/2010 15:57	2.1	12799	Japan
4	4	12799 2010-06-28 15:57:00	4				
513796	22074	6 RIBBONS SHIMMER...	12	6/28/2010 15:57	1.65	12799	Japan
4	4	12799 2010-06-28 15:57:00	4				
513796	85032B	BLOSSOM IMAGES GI...	6	6/28/2010 15:57	2.1	12799	Japan
4	4	12799 2010-06-28 15:57:00	4				
513796	22078	RIBBON REEL LACE ...	10	6/28/2010 15:57	2.1	12799	Japan
4	4	12799 2010-06-28 15:57:00	4				
513796	21259	VICTORIAN SEWING ...	2	6/28/2010 15:57	5.95	12799	Japan
4	4	12799 2010-06-28 15:57:00	4				
513796	85017C	ENVELOPE 50 CURIO...	12	6/28/2010 15:57	0.85	12799	Japan
4	4	12799 2010-06-28 15:57:00	4				
513796	85178	VICTORIAN SEWING KIT	12	6/28/2010 15:57	1.25	12799	Japan
4	4	12799 2010-06-28 15:57:00	4				
513796	85176	SEWING SUSAN 21 N...	12	6/28/2010 15:57	0.85	12799	Japan
4	4	12799 2010-06-28 15:57:00	4				
513796	85049D	BRIGHT BLUES RIBB...	12	6/28/2010 15:57	1.25	12799	Japan
4	4	12799 2010-06-28 15:57:00	4				
513796	22081	RIBBON REEL FLORA...	10	6/28/2010 15:57	1.65	12799	Japan
4	4	12799 2010-06-28 15:57:00	4				
566488	22600	CHRISTMAS RETROSP...	12	9/13/2011 10:16	0.85	12940	United Kingdom
4	3	12940 2011-10-24 14:04:00	4				
571270	22696	WICKER WREATH LARGE	2	10/16/2011 12:09	1.95	12940	United Kingdom
4	3	12940 2011-10-24 14:04:00	4				
571270	21619	4 VANILLA BOTANIC...	3	10/16/2011 12:09	1.25	12940	United Kingdom
4	3	12940 2011-10-24 14:04:00	4				
566488	23174	REGENCY SUGAR BOW...	4	9/13/2011 10:16	4.15	12940	United Kingdom
4	3	12940 2011-10-24 14:04:00	4				
571270	23333	IVORY WICKER HEAR...	4	10/16/2011 12:09	1.25	12940	United Kingdom
4	3	12940 2011-10-24 14:04:00	4				
+-----+-----+-----+-----+-----+-----+-----+-----+							
-+-----+-----+-----+-----+							

only showing top 20 rows

```
In [46]: recency_df.columns
```

```
Out[46]: ['Invoice',  
          'StockCode',  
          'Description',  
          'Quantity',  
          'InvoiceDate',  
          'Price',  
          'Customer_ID',  
          'Country',  
          'Monetary',  
          'Frequency',  
          'Customer ID',  
          'first(InvoiceDate)',  
          'Recency']
```

```
In [47]: recency_df = recency_df.drop('first(InvoiceDate)', 'Customer ID')
```

```
In [48]: RFM_matrix = recency_df.select(['Customer_ID', 'Recency', 'Frequency', 'Monetary'])  
RFM_matrix.show()
```

```
[Stage 75:=====> (194 + 4) / 200]
```

Customer_ID	Recency	Frequency	Monetary
12799	4	4	4
12799	4	4	4
12799	4	4	4
12799	4	4	4
12799	4	4	4
12799	4	4	4
12799	4	4	4
12799	4	4	4
12799	4	4	4
12799	4	4	4
12799	4	4	4
12799	4	4	4
12799	4	4	4
12799	4	4	4
12799	4	4	4
12799	4	4	4
12940	4	3	4
12940	4	3	4
12940	4	3	4
12940	4	3	4
12940	4	3	4

only showing top 20 rows

```
In [49]:
```

```
RFM_matrix.count()
```

```
Out[49]:
```

```
797885
```

Find the number of customers in each of the 6 categories in the table above

"Best customers" segment with RFM 111

```
In [50]: RFM_matrix.filter((RFM_matrix['Recency'] == 1) &
                             (RFM_matrix['Frequency'] == 1) &
                             (RFM_matrix['Monetary'] == 1)).show()
```

```
+-----+-----+-----+-----+
|Customer_ID|Recency|Frequency|Monetary|
+-----+-----+-----+-----+
+-----+-----+-----+-----+
```

"Loyal customers" segment with RFM X1X

```
In [51]: RFM_matrix.filter( (RFM_matrix['Frequency'] == 1)).show()
```

```
[Stage 89:=====>(197 + 3) / 200]
```

Customer_ID	Recency	Frequency	Monetary
13623	4	1	4
13623	4	1	4
13623	4	1	4
13623	4	1	4
13623	4	1	4
13623	4	1	4
13623	4	1	4
13623	4	1	4
13623	4	1	4
13623	4	1	4
13623	4	1	4
13623	4	1	4
13623	4	1	4
13623	4	1	4
13623	4	1	4
13623	4	1	4
13623	4	1	4
13623	4	1	4
13623	4	1	4
13623	4	1	4

only showing top 20 rows

"Big spenders" segment with RFM XX1

```
In [52]: RFM_matrix.filter((RFM_matrix['Monetary'] == 1)).show()
```

Customer_ID	Recency	Frequency	Monetary
12757	4	1	1
12757	4	1	1
12757	4	1	1
12757	4	1	1
15202	4	2	1
15202	4	2	1
15202	4	2	1
15202	4	2	1
15202	4	2	1
15202	4	2	1
15202	4	2	1
15202	4	2	1
14096	4	1	1
14096	4	1	1
14096	4	1	1
14096	4	1	1
14096	4	1	1
14096	4	1	1
14096	4	1	1
14096	4	1	1
14096	4	1	1
14096	4	1	1

only showing top 20 rows

"Almost lost" segment with RFM 311

```
In [53]: RFM_matrix.filter((RFM_matrix['Recency'] == 3) &
                        (RFM_matrix['Frequency'] == 1) &
                        (RFM_matrix['Monetary'] == 1)).show()
```

Customer_ID	Recency	Frequency	Monetary

"Lost customers" segment with RFM 411



[Stage 146:=====> (194 + 4) / 200]

Customer_ID	Recency	Frequency	Monetary
12799	4	4	4
12799	4	4	4
12799	4	4	4
12799	4	4	4
12799	4	4	4
12799	4	4	4
12799	4	4	4
12799	4	4	4
12799	4	4	4
12799	4	4	4
12799	4	4	4
12799	4	4	4
12799	4	4	4
12799	4	4	4
12799	4	4	4
13289	4	4	4
13289	4	4	4
13289	4	4	4
13289	4	4	4
13289	4	4	4

only showing top 20 rows